How much do loop reordering and cache blocking speed up matrix multiplication, and when does a sparse (CSR) representation pay off?

Big Data - Individual Assignment 2: Optimised matrix multiplication approaches and sparse matrices
Grado en Ciencia e Ingeniería de Datos, Universidad de Las Palmas de Gran Canaria, 2026-2027
Author: Guillermo Oltra Rodriguez - **Date:** October 2026
Materials: GitHub repository https://github.com/shisher06/bigdata-matmul-assignment2 (code, input configurations, raw measurements, analysis scripts; see `README.md` and Section 8). It reuses the validated C implementation and the input files of Assignment 1.

Abstract. This study continues Assignment 1 (naive `i-j-k` matrix multiplication in C, Python and Java) with two questions. (A) How much do two dense optimisations, loop reordering (`i-k-j`) and cache blocking (64 x 64 tiles), speed up the Assignment 1 C kernel, and do they remove the cache effects found there? (B) For which densities is an own sparse CSR x CSR multiplication (Gustavson's algorithm) faster than the best dense method? All methods are in C (gcc 11.2 `-O2`, one thread), use the same input files (Part A reuses the Assignment 1 files), give exactly the same result bit for bit, and were measured on one laptop (Ryzen 7 4800H, Windows 11) with a new process per run and 5 repetitions. Loop reordering removed the slowdowns at n = 512, 1024, 1536 and 2048 and the growth faster than n^3 (log-log slope 3.05 instead of 3.74); blocking was the fastest dense method (about 0.33 ns per inner iteration up to n = 4000; 5.7 times faster than the baseline at n = 2000 and 15 times at n = 3000). Blocking was also faster than `i-k-j` at small n, where tiling cannot matter, so its gain over `i-k-j` is only explained by cache reuse for large n. CSR was faster than blocking up to 50 % nonzeros (from about 10 000 times at 0.1 % to 1.13 times at 50 %) and slower above a crossover of about 53 % for n = 1000 and 2000, in agreement with a simple work model; the conversion to CSR (up to 100 times the product time at 0.1 %) and the fill-in of the result (almost dense from 5 %) reduce the advantage. The conclusions only apply to these implementations, uniformly random sparse inputs and this machine.

1. Introduction

In Assignment 1 I compared the basic triple-loop multiplication (`i-j-k`) in Python, Java and C. The main open point of that study was the memory hierarchy: in C the time grew faster than n^3 (log-log slopes of 3.2 to 3.6), it jumped when matrix `B` stopped fitting in the 4 MiB L3 cache (between n = 600 and 800), and at n = 512, 1024 and 1536 it was 6 to 7 times slower than at nearby sizes. I explained these effects by the column-wise reading of `B` in the inner loop, but I could not test the explanation. This assignment tests it by changing *only* the way the loops walk through memory, and then looks at a different way to save work: not multiplying zeros at all.

Research questions.

- **Part A (dense optimisation).** How much do two memory-oriented optimisations, *loop reordering* (`i-k-j`) and *cache blocking* (tiled `i-k-j`), reduce the runtime of the Assignment 1 C triple loop, and do they remove the effects of the memory hierarchy found there?
- **Part B (sparse matrices).** For which tested densities and sizes is a sparse multiplication (CSR x CSR, Gustavson's algorithm) faster than the best dense method on the same matrices, and how much of that advantage is lost to the conversion to CSR and to the fill-in of the result?

- **What changes:** the method (`ijk`, `ikj`, `blocked`, `csr`), the size `n` and, in Part B, the density (percentage of nonzero entries) of the inputs.
- **What is measured:** time per multiplication (median and IQR of 5 repetitions), peak memory of the process and, in Part B, the conversion time, the memory of the CSR structures and the number of nonzeros of the inputs and of the result.
- **What stays the same:** the language (C), the compiler and its flags, the numeric type (`float64`), the timing boundary, the machine and, within each comparison, the input files.

Hypotheses (written before the main campaigns, from the Assignment 1 results and a short preliminary run at n = 512 and 1000):

- **H1.** `ikj` will be faster than `ijk` for large `n` and will remove the slowdowns at powers of two, because its inner loop reads `B` and `C` row by row instead of reading `B` by columns.
- **H2.** `blocked` will be faster than `ikj` once the matrices no longer fit in the cache (n of the order of 700 or more), and close to it for small `n`. With blocking the time should follow n^3 closely (constant cost per iteration).
- **H3.** The three dense methods will use the same memory (3 x 8 x n^2 bytes plus the C runtime), since they only change the order of the loops.
- **H4.** CSR will be much faster than the dense methods at low densities, because its work is proportional to the number of useful products (about n^3d^2 for density d) instead of n^3; the advantage will shrink as the density grows, and above some density (expected between 10 % and 50 %) dense `blocked` will be faster, because each sparse product costs more than a dense one (indirect indices, scattered accesses to the accumulator).
- **H5.** The result `C` will be much denser than `A` and `B` (fill-in), so the output part of the CSR algorithm and its memory will grow faster than the density of the inputs.

2. Background

The three dense loop orders. `C = A x B` needs n^3 multiply-adds whatever the loop order [1]. The order only changes *which memory positions are read one after another*. Matrices are stored row by row (row-major), and the CPU reads memory in 64-byte cache lines (8 doubles) [2]:

- `i-j-k` (Assignment 1): the inner loop runs over `k` and reads `B[k][j]`, a **column** of `B`; two consecutive reads are 8n bytes apart, so each one uses a different cache line and only 8 of its 64 bytes. When `B` ($8n^2$ bytes) is larger than the cache, these lines must come again from main memory for every row of `C`.
- `i-k-j` (loop reordering): the inner loop runs over `j` and computes `C[i][j] += a * B[k][j]` with a fixed `a = A[i][k]`. Both `B` and `C` are read along a **row**, so every cache line brings 8 useful values, and the hardware prefetcher can predict the next ones. The amount of arithmetic is the same.
- **Cache blocking (tiling)**: the `i-k-j` loops are run on square tiles of bs x bs elements, so that a tile of `B` and the corresponding part of the rows of `C` are reused many times while they are still in a fast cache, instead of streaming the whole of `B` ($8n^2$ bytes) from memory for every row of `C` [1], [3]. Lam, Rothberg and Wolf showed that the best block size depends on the cache size and on conflicts between addresses, and that the block should usually be smaller than the cache [3].

Cache conflicts at powers of two. In a set-associative cache each address can only go to a few places ("ways") of one set, chosen by some bits of the address [2]. If consecutive accesses are 8n bytes apart and 8n is a multiple of 4096 (n = 512, 1024, 1536, 2048), they all fall in the same few sets and evict each other, so the cache behaves as if it were much smaller. This is the explanation of the Assignment 1 slowdowns; in `ijk` it affects the column walk of `B`, which `ikj` and `blocked` no longer do.

Sparse matrices and CSR. A matrix is *sparse* when most of its entries are zero. Here the **density** `d` is the percentage of nonzero entries and the **sparsity** is `100 - d`. The *Compressed Sparse Row* (CSR) format stores, row by row, only the nonzero values (`val`), their column indices (`col`) and, for each row, where it starts (`row_ptr`) [4]. With 8-byte values, 4-byte column indices and 8-byte row pointers, a matrix takes `12 nnz + 8(n + 1)` bytes instead of `8n²`.

Gustavson's algorithm [5] computes `C = A x B` row by row: row `i` of `C` is the sum of the rows `k` of `B` multiplied by the nonzeros `A[i][k]`. Only nonzero pairs are multiplied, so the work is proportional to the number of useful products. If the nonzeros are placed independently with probability `d`, the expected number of useful products is n^3d^2, and an entry of `C` is nonzero with probability `1 - (1 - d²)^n` (it is zero only if none of the n products is nonzero). For example with n = 1000 and d = 5 %, `C` is about 92 % nonzero although the inputs are 95 % zeros: the output becomes dense very quickly.

Hardware. The Ryzen 7 4800H (Zen 2) has, per core, 32 KiB of L1 data cache (8-way) and 512 KiB of L2 (8-way), and each group of 4 cores shares 4 MiB of L3 (16-way) [6]. A 64 x 64 tile of doubles takes 32 KiB, so three tiles (of `A`, `B` and `C`) fit in L2.

3. Methodology

3.1 What is reused from Assignment 1

The Assignment 1 materials (repository https://github.com/shisher06/bigdata-matmul-assignment1, commit `e947d18`, folder `practica1/`) were copied into `practica2/` and extended, so that this submission is self-contained:

Reused	How
C worker `src/c/matmul.c`	The kernel `matmul_ijk` is **unchanged** and is the baseline. Timer, file reading, warm-up, batching, checksum and memory reading are the same code. New: `--method`, `--block`, the two optimised kernels and the CSR code.
Input files and generator	`data/A_n.bin`, `data/B_n.bin` are **the same files** (same generator, seed 2026, same SHA-256, checked by a test). New: a sparse generator.
Verifier and tolerances	`verify.py` without changes (NumPy, `rtol = 1e-9`, `atol = 1e-12`).
Benchmark runner	Same protocol (Section 3.5); it compares methods instead of languages.
Build, environment, analysis, PDF	Same scripts, adapted (only C `-O2`; new tables and plots).

Only C is used. The Assignment 1 study showed that the memory effects that the optimisations target are clearly visible in C, while in Python the interpreter cost (about 70 ns per iteration) hides them.

3.2 Implementations

All four kernels are in the same binary, compiled with the same flags, and are selected with `--method`. No matrix library is used inside the timed code.

```
/* ijk - baseline, Assignment 1 (unchanged)      */  /* ikj - loop reordering                    */
for (i = 0; i < n; i++)                                 for (i = 0; i < n; i++)
  for (j = 0; j < n; j++)                                 for (k = 0; k < n; k++) {
    for (k = 0; k < n; k++)                                 a = A[i][k];
      C[i][j] += A[i][k] * B[k][j];                         for (j = 0; j < n; j++)
                                                              C[i][j] += a * B[k][j];
                                                            }
/* blocked - ikj on bs x bs tiles (bs = 64) */
for (ii = 0; ii < n; ii += bs)
  for (kk = 0; kk < n; kk += bs)
    for (jj = 0; jj < n; jj += bs)
      for (i = ii; i < min(ii+bs, n); i++)
        for (k = kk; k < min(kk+bs, n); k++) {
          a = A[i][k];
          for (j = jj; j < min(jj+bs, n); j++) C[i][j] += a * B[k][j];
        }
```

The CSR kernel (`matmul_csr`) follows Gustavson's algorithm [5]. For each row `i` it walks the nonzeros `A[i][k]` (in increasing `k`) and, for each, the nonzeros `B[k][j]`, adding `A[i][k] * B[k][j]` to a dense accumulator row `acc[j]`. An array `mark[j]` records whether column `j` already appeared in row `i` (then `acc[j]` is valid; otherwise it is set to 0.0 and `j` is added to a list). At the end of the row the columns are sorted (with `qsort` if there are few, or by scanning `mark` if there are more than n/8) and copied into the CSR arrays of `C`, which grow by doubling. The conversion from the dense input to CSR (`csr_from_dense`) counts the nonzeros and then fills the three arrays.

Method	Storage	Work	What changes compared with `ijk`
`ijk`	dense, $8n^2$ B per matrix	n^3	-
`ikj`	dense	n^3	loop order: inner loop along rows of `B` and `C`
`blocked`	dense	n^3	`ikj` on 64 x 64 tiles (block chosen in Section 4.3)
`csr`	CSR, 12 nnz + 8(n + 1) B per matrix	about n^3d^2 products + O(n) per row	only nonzero products; indirect indexing

Same result, bit for bit. In all four kernels each `C[i][j]` starts at 0.0 and receives the products in increasing order of `k`, without FMA (`-ffp-contract=off`): in `blocked` the `k` tiles are visited in increasing order, and in `csr` the skipped products are exact zeros, and adding `+0.0` does not change a `double`. So all methods must produce **exactly the same bits** as the Assignment 1 baseline, which gives a very strict correctness check (Section 3.7).

No SIMD. With gcc 11.2, `-O2` does not enable the automatic vectorisation (`gcc -Q --help=optimizers -O2` lists `-ftree-loop-vectorize` as disabled; it is enabled at `-O3`) [7]. So the differences between methods come from memory access and loop overhead, not from vector instructions, which are left for Assignment 3.

3.3 Inputs

- **Part A** uses the Assignment 1 files `data/A_n.bin` and `data/B_n.bin` (values uniform in [0, 1), `float64`, row by row, CPython `random` with base seed 2026), plus the same generator for the new sizes. A test checks that the generator still reproduces the SHA-256 of the Assignment 1 manifest.
- **Part B** uses new files `data_sparse/A_<n>_d<pct>.bin`, written in the **same dense format** (zeros included) so that the dense and sparse methods read exactly the same logical matrices. For each cell, a uniform number `u` is drawn; if `u < d` the cell gets a value `1 - random()` in (0, 1], otherwise 0. So the nonzeros are placed **uniformly and independently** (no structure), with the declared density on

average. The seed of each matrix depends on n and d (rule in `README.md`), and `data_sparse/manifest.json` stores the seed, the real number of nonzeros and the SHA-256 of each file.

- **Densities:** 0.1, 0.5, 1, 2, 5, 10, 20, 30 and 50 % (sparsity 99.9 % to 50 %).

3.4 Environment

Item	Value
Machine	Laptop, AMD Ryzen 7 4800H, 8 cores / 16 threads, 2.9 GHz (boost up to 4.2 GHz)
Caches [6]	L1d 32 KiB and L2 512 KiB per core; L3 2 x 4 MiB
Memory	15.4 GiB
Operating system	Windows 11 Home (build 26300)
Power	Plugged in, default power plan
C	gcc 11.2.0 (MinGW-w64), `-std=c11 -O2 -ffp-contract=off`, the four methods in one binary
Threads	One computing thread; `OMP/OPENBLAS/MKL_NUM_THREADS=1` in the measured processes
Python (runner, analysis)	CPython 3.11.9, NumPy 2.4.6 (only for checking), matplotlib 3.11.2, psutil 7.2.2

This information is saved automatically in `results/*/environment.md` and `build/build_info.json`.

3.5 Timer and measurement protocol (same as Assignment 1)

- **What the timer includes:** only the kernel call: allocating `C`, initialising it (or building the CSR of `C`) and running the loops. Reading the files, the warm-up, the checksum and saving the result are outside. Timer: `QueryPerformanceCounter` (monotonic). The result is always used afterwards (a checksum of all elements), so the compiler cannot skip the computation.
- **Conversion (Part B):** for `csr`, the conversion of `A` and `B` from the dense input to CSR is timed **separately** (once per process). Results are given both for the product alone and for product + conversion; the output stays in CSR (it is not converted back to dense inside any timer).
- **A new process for each repetition** and method; **interleaved order** (for each case and repetition all methods run, and the order rotates in each repetition).
- **5 repetitions** per configuration (3 if the first one takes more than 120 s).
- **Short tasks:** if one multiplication takes less than 100 ms, a batch is timed together and the time is divided by the batch size. For `csr` the batch size is always calibrated with one extra (discarded) run, because its cost is not cubic.
- **Warm-up:** one run on the 64 x 64 corner of the matrices, not included in the results.
- **Statistics:** median and inter-quartile range (IQR); all raw measurements are kept in `raw.csv`.

3.6 Memory metrics

- **Peak working set of the process** (the Assignment 1 metric): read by the runner with `GetProcessMemoryInfo` when the process ends. For `csr` it includes reading the dense input files, so it does not show the saving of CSR by itself.
- **Working set after the product** (new): the current working set of the process right after the timed region. In `csr` the dense inputs have already been freed, so it reflects the CSR structures plus the C runtime (about 4 MB).
- **Data-structure model**: exact bytes of the structures used by each method: 3 x $8n^2$ for the dense methods, `12 nnz + 8(n + 1)` per matrix (A, B and C) for CSR.

3.7 Correctness

Before measuring, **24 automated tests** (`tests/`) were run with every method: small cases computed by hand (1x1, 2x2, 3x3), identity and zero matrices, non-commutativity, exact integer results, comparison with NumPy, **bit identity with the Assignment 1 Python kernel** (n = 40 and 67, which is not a multiple of any block size; `blocked` is also tested with block 5), sparse inputs (0 %, 100 % and random densities, with the nonzero counts), the CSR memory model, batching and invalid inputs; plus tests of the input files (Assignment 1 hashes), the sparse generator and the verifier.

During each campaign, as in Assignment 1, in repetition 1 of each method and case the result is saved after the timer and compared with NumPy's `A @ B` (`rtol = 1e-9`, `atol = 1e-12`; NumPy adds in a different order, so differences around 1e-15 are expected). The SHA-256 of the result must be the same for all methods, and in repetitions 2 to 5 the checksum must be equal to repetition 1.

3.8 Sizes and limits (decided before measuring)

- **Part A:** n = 10, 25, 50, 100, 200, 300, 400, 500, **512**, 600, 800, 1000, **1024**, 1200, 1500, **1536**, 2000, **2048**, 2500, 3000, 4000: the Assignment 1 sizes, the powers of two (and 1536) where Assignment 1 found slowdowns, and larger sizes to find the limits of each method.
- **Block size (separate experiment):** bs = 16, 32, 64, 128, 256 at n = 500, 512, 1000, 1024, 2000.
- **Part B:** n = 1000 and 2000 with the nine densities and the methods `csr`, `blocked` and `ijk`; n = 4000 with 0.1 % and 1 % (`csr` and `blocked`; `ijk` at n = 4000 is measured in Part A).
- **Limits (as in Assignment 1):** 300 s per process and 4096 MB per process. A run over a limit is recorded as a result with no time, and that method does not try larger cases.

3.9 Separate experiments

Two experiments are kept apart from the main comparisons (own folders, not mixed in the main tables):

- `02_A_extra_bloques`: the block-size sweep (Section 4.3), run first to choose the block.
- `05_B_extra_densidades_altas`: CSR and `blocked` at densities of 60, 70, 85 and 100 % (n = 1000 and 2000). **It was added after the main Part B campaign**, because CSR was still faster at the highest declared density (50 %), so the crossover had not been found.

4. Results

Data: `results/01_A_densa/` (Part A: 307 runs, 306 correct and 1 timeout), `02_A_extra_bloques/` (125 runs), `03_B_dispersa/` (270 runs), `04_B_dispersa_n4000/` (20 runs), `05_B_extra_densidades_altas/` (80 runs). Each folder has `raw.csv`, `summary.md` (complete tables), `campaign.log`, `environment.md` and `plots/`. All campaigns were run on 27 Oct 2026 in one session with the laptop plugged in; Part A took 41 minutes and Part B 26 minutes.

4.1 Correctness

- The 24 tests passed before every campaign.
- **All methods produced exactly the same result** (same SHA-256 of `C`) in every comparable case: the 21 sizes of Part A, the 26 (n, density) cases of Part B and the cases of the separate experiments. So `ikj`, `blocked` and `csr` compute bit for bit the same numbers as the Assignment 1 baseline.
- Maximum relative error against NumPy: 8.3e-15 (tolerance 1e-9). In repetitions 2 to 5 the checksum was always equal to repetition 1.

4.2 Part A: runtime of the three dense methods

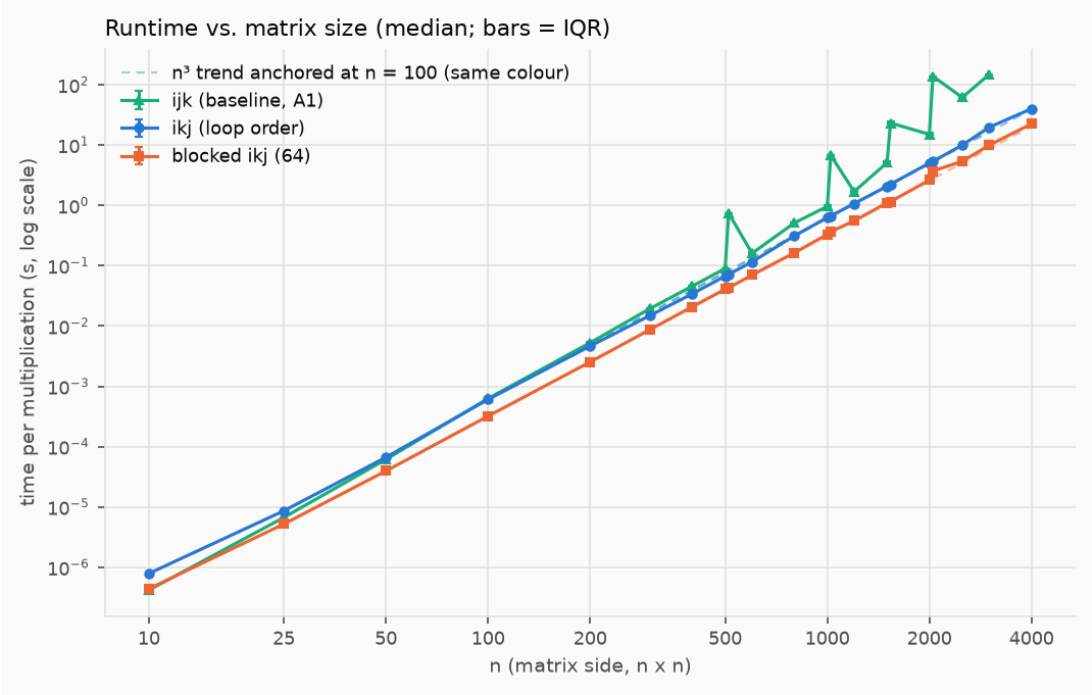


Figure 1. Time per multiplication against n (log-log), median of 5 repetitions (3 for ijk at n = 2048 and 3000); the bars (IQR) are mostly smaller than the markers.

n	`ijk` (baseline)	`ikj`	`blocked` (64)	speedup `ikj`	speedup `blocked`
10	0.418 µs	0.788 µs	0.439 µs	0.53x	0.95x
50	61.5 µs	66.5 µs	39.5 µs	0.92x	1.56x
100	627 µs	605 µs	320 µs	1.04x	1.96x
200	5.28 ms	4.64 ms	2.53 ms	1.14x	2.09x
500	89.9 ms	67.0 ms	40.7 ms	1.34x	2.21x
512	**739 ms**	71.9 ms	43.0 ms	**10.3x**	**17.2x**
800	511 ms	312 ms	164 ms	1.64x	3.13x
1000	960 ms	627 ms	329 ms	1.53x	2.92x
1024	**6.75 s**	667 ms	366 ms	**10.1x**	**18.4x**
1500	5.09 s	2.07 s	1.10 s	2.46x	4.65x
1536	**23.2 s**	2.24 s	1.15 s	**10.4x**	**20.1x**
2000	15.1 s	5.02 s	2.63 s	2.99x	5.73x
2048	**138 s**	5.36 s	3.70 s	**25.8x**	**37.3x**
2500	61.5 s	9.97 s	5.38 s	6.17x	11.4x
3000	148 s	19.7 s	9.92 s	7.52x	14.9x
4000	timeout (> 300 s)	40.0 s	22.5 s	-	-

Speedup = median time of `ijk` / median time of the method. Complete table (21 sizes, with Q1, Q3 and IQR): `results/01_A_densa/summary.md`. Variability was low: IQR below 2 % of the median in 32 of the 62 configurations and below 5 % in 54; the largest was 11 % (`ikj`, n = 3000).

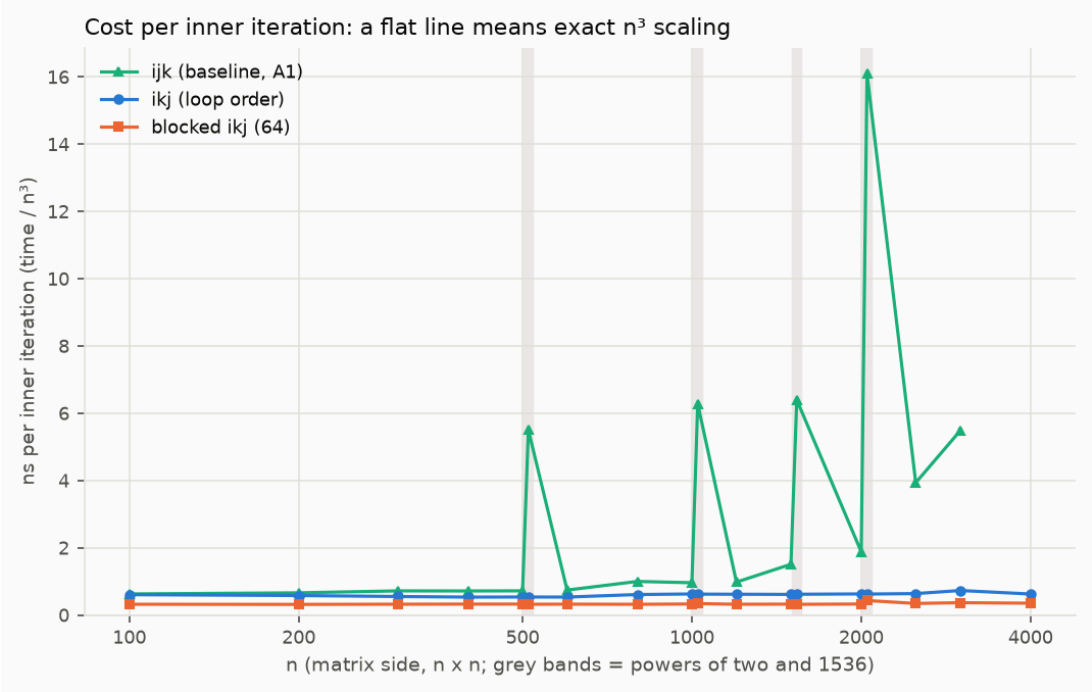


Figure 2. Time divided by n^3 (ns per inner iteration). A flat line means that the time grows exactly as n^3. Grey bands: n = 512, 1024, 1536 and 2048.

- **Baseline `ijk`.** It reproduces the Assignment 1 behaviour: about 0.42 to 0.74 ns per iteration up to n = 600, about 1 ns between 800 and 1200, and then a fast increase (1.5 ns at 1500, 1.9 at 2000, 3.9 at 2500, 5.5 at 3000), with peaks of 5.5 to 6.4 ns at 512, 1024 and 1536 and **16.1 ns at 2048**. The log-log slope for n >= 100 is **3.74** (3.58 without the powers of two). At n = 4000 the first run passed the 300 s limit (recorded as a timeout), so its largest size within the limits is 3000.
- **`ikj`.** Between 0.53 and 0.73 ns per iteration from n = 25 to 4000, **without any peak at the powers of two**. Slope **3.05**.
- **`blocked`.** Between 0.32 and 0.37 ns per iteration from n = 25 to 4000 (0.43 at n = 2048), about 1.9 times less than `ikj` at every size. Slope **3.04**.
- **Small matrices.** At n = 10 `ikj` is 1.9 times *slower* than `ijk`, and at n = 25 to 50 it is still slightly slower; from n = 100 it is faster.
- **Comparison with Assignment 1.** The `ijk` medians are within -5 % and +11 % of the Assignment 1 values up to n = 1500 (same code, flags, inputs and machine, measured 13 days later), but at n = 2000 the new median is 15.1 s against 10.8 s (1.39 times).

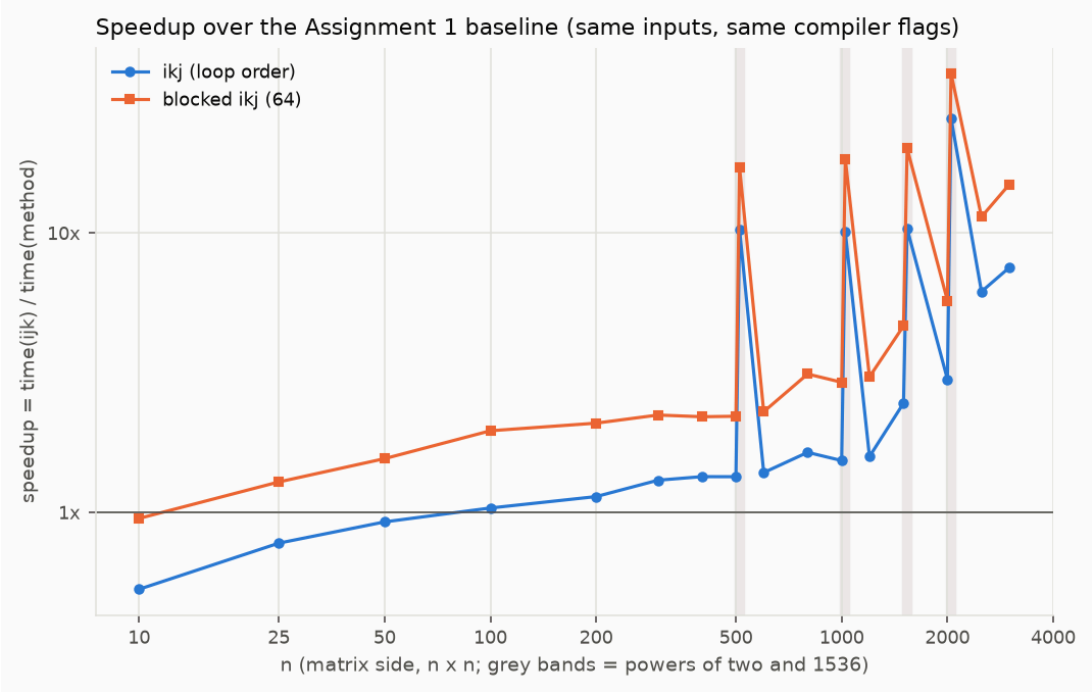


Figure 3. Speedup of the optimised methods over the Assignment 1 baseline.

4.3 Separate experiment: block size

n	block 16	block 32	block 64	block 128	block 256
500	0.364	**0.333**	0.338	0.436	0.441
512	0.391	0.345	**0.324**	0.421	0.415
1000	0.356	0.334	**0.327**	0.413	0.419
1024	0.387	**0.335**	0.344	0.472	0.417
2000	0.355	0.330	**0.329**	0.413	0.422

ns per inner iteration (median of 5); best block in bold. Blocks of 32 and 64 are close (within 7 %), 16 is 8 to 21 % slower than the best and 128 or 256 are 24 to 41 % slower. The value 64 used in the main campaign is the best or within 3 % of the best at every size. At n = 500 the first two repetitions of blocks 32 and 64 were about twice as slow as the other three (the laptop was also being used at the start of the session); the medians are not affected, and no measurement was removed.

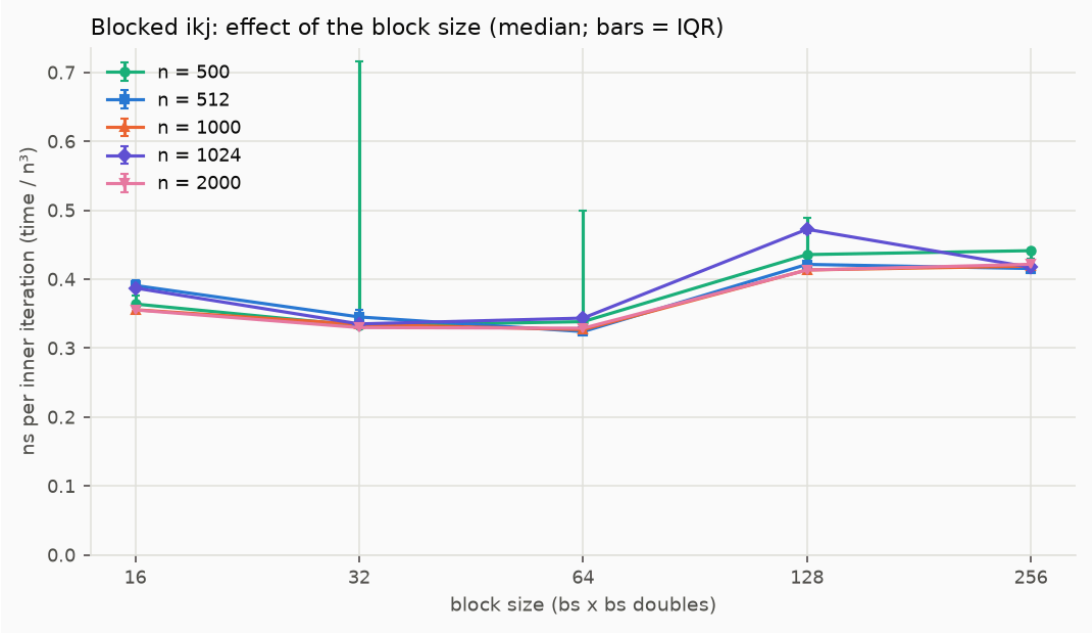


Figure 4. Blocked i-k-j: cost per iteration against the block size.

4.4 Part A: memory

The peak working set is the same for the three dense methods (for example 28.5 MB at n = 1000, 100.5 MB at n = 2000 and 388.5 MB at n = 4000) and follows the storage model 3 x 8 x n^2 (24, 96 and 384 MB) plus about 4.5 MB of the C runtime. The only differences (up to 2.8 MB, at n = 512 and 600) come from the batches used for short runs. Largest peak: 388.5 MB, far from the 4096 MB budget.

4.5 Part B: sparse inputs, CSR against the dense methods

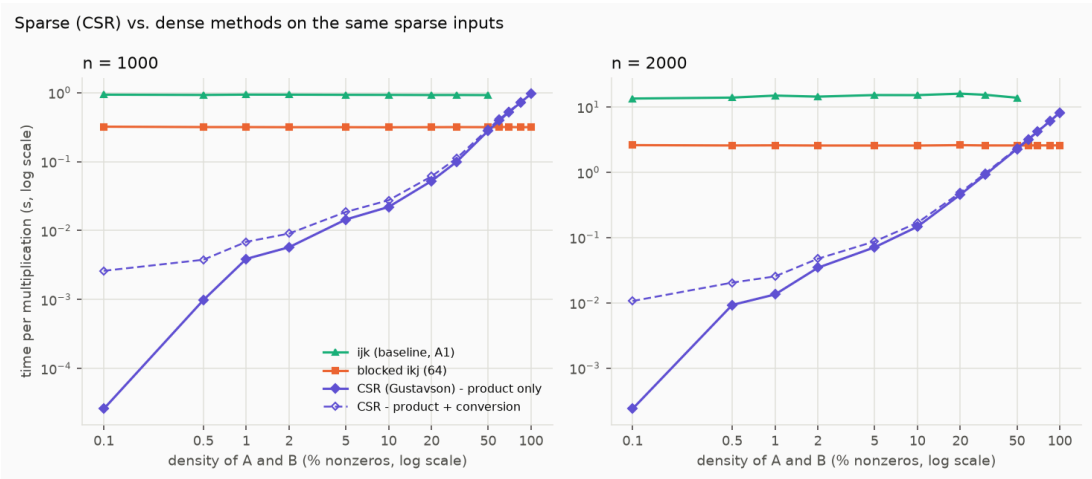


Figure 5. Time per multiplication against the density of A and B, with the same input files for all methods (main campaign and the high-density experiment). Dashed: CSR product plus the conversion of A and B to CSR.

n	density	nnz(A)	density of C	CSR product	conversion	`blocked`	`ijk`	`blocked` / CSR	`blocked` / (CSR + conv.)
1000	0.1 %	969	0.11 %	26.1 µs	2.58 ms	321 ms	934 ms	12 300x	123x
1000	1 %	9 964	9.4 %	3.88 ms	2.98 ms	316 ms	935 ms	81x	46x
1000	5 %	49 850	91.8 %	14.5 ms	4.15 ms	315 ms	929 ms	22x	17x

n	density	nnz(A)	density of C	CSR product	conversion	`blocked`	`ijk`	`blocked` / CSR	`blocked` / (CSR + conv.)
1000	20 %	199 650	100 %	52.3 ms	8.73 ms	315 ms	923 ms	6.0x	5.2x
1000	50 %	500 386	100 %	280 ms	17.3 ms	316 ms	921 ms	1.13x	1.06x
1000	60 %	600 284	100 %	396 ms	16.0 ms	317 ms	-	0.80x	0.77x
1000	100 %	1 000 000	100 %	960 ms	5.63 ms	316 ms	-	0.33x	0.33x
2000	0.1 %	4 035	0.20 %	238 µs	10.4 ms	2.61 s	13.5 s	11 000x	245x
2000	1 %	39 982	18.1 %	13.4 ms	12.0 ms	2.59 s	15.0 s	193x	102x
2000	5 %	200 753	99.3 %	70.9 ms	16.3 ms	2.57 s	15.3 s	36x	30x
2000	20 %	800 612	100 %	456 ms	34.1 ms	2.61 s	16.1 s	5.7x	5.3x
2000	50 %	1 999 895	100 %	2.28 s	69.5 ms	2.58 s	13.9 s	1.13x	1.10x
2000	60 %	2 399 891	100 %	3.17 s	64.3 ms	2.58 s	-	0.82x	0.80x
2000	100 %	4 000 000	100 %	8.23 s	22.4 ms	2.59 s	-	0.31x	0.31x
4000	0.1 %	16 029	0.40 %	2.14 ms	40.8 ms	21.9 s	-	10 200x	509x
4000	1 %	159 986	32.9 %	91.8 ms	46.3 ms	21.6 s	-	236x	157x

Medians of 5 repetitions; all densities in `results/03_B_dispersa/summary.md`, `results/05_B_extra_densidades_altas/summary.md` and `results/04_B_dispersa_n4000/summary.md`.

- **CSR is faster than the best dense method at every density up to 50 %**, by more than 10 000 times at 0.1 % (product only) and by 1.13 times at 50 %. It is slower from 60 % on, and 3.0 to 3.2 times slower at 100 %. Interpolating in log-log scale, the crossover is at **53.3 % (n = 1000) and 53.6 % (n = 2000)** for the product alone, and at 51.8 % and 52.8 % including the conversion.
- **The dense methods do not depend on the density**: `blocked` takes 314 to 321 ms (n = 1000) and 2.57 to 2.61 s (n = 2000) at every density, the same as with the dense inputs of Part A.
- **Conversion.** It takes 2.6 to 17 ms at n = 1000 and 10 to 70 ms at n = 2000. At 0.1 % it is 100 times longer than the CSR product, so including it reduces the advantage from about 12 000x to 123x (n = 1000) and 245x (n = 2000). It is largest at 50 to 60 % and *decreases* again towards 100 % (at n = 2000: 10.4 ms at 0.1 %, 69.5 ms at 50 %, 22.4 ms at 100 %).
- **Larger size.** At n = 4000 and 0.1 % density, CSR needs 2.1 ms (43 ms with the conversion) against 21.9 s for `blocked`.

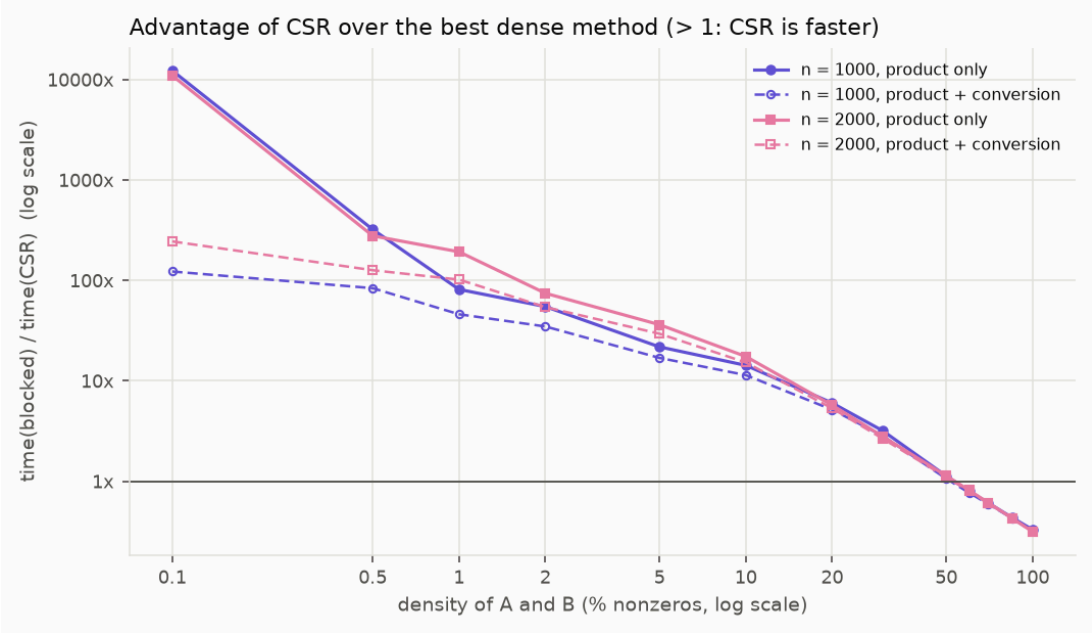


Figure 6. Time of blocked divided by time of CSR (above 1: CSR is faster).

4.6 Part B: fill-in of the result and cost per product

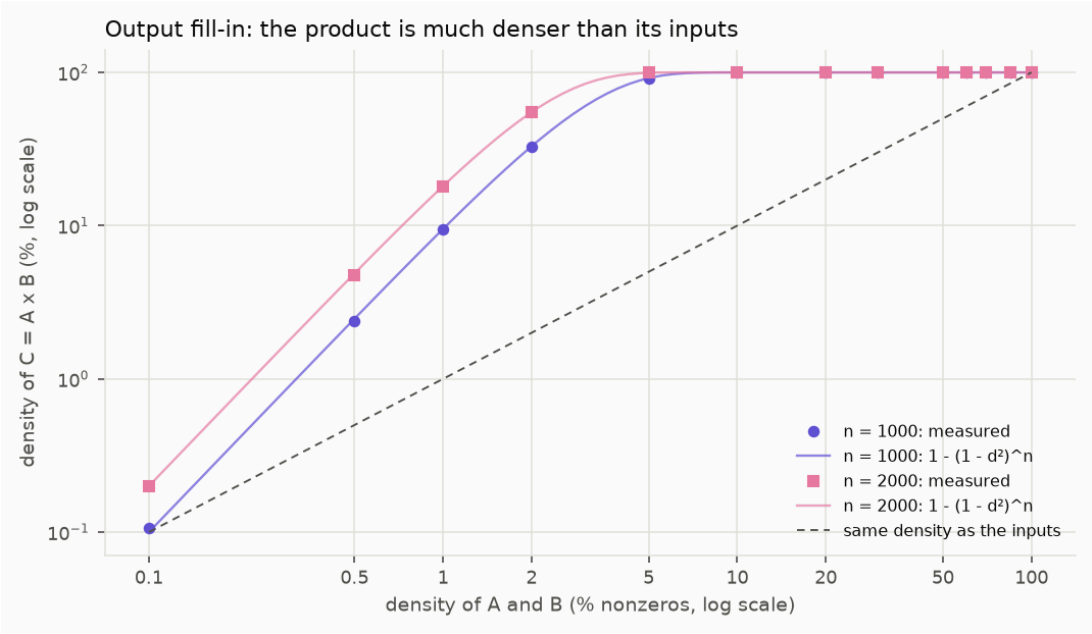


Figure 7. Density of C = A x B against the density of the inputs; lines: the model 1 - (1 - d^2)^n for independent uniform nonzeros.

The measured density of `C` agrees with the model 1 - (1 - d^2)^n at every point (for example 9.45 % measured and expected at n = 1000, d = 1 %; 55.0 % at n = 2000, d = 2 %). With inputs that are 99 % zeros, `C` is already 9.4 % (n = 1000) or 18.1 % (n = 2000) nonzero, and from d = 5 % it is almost completely dense (91.8 % and 99.3 %).

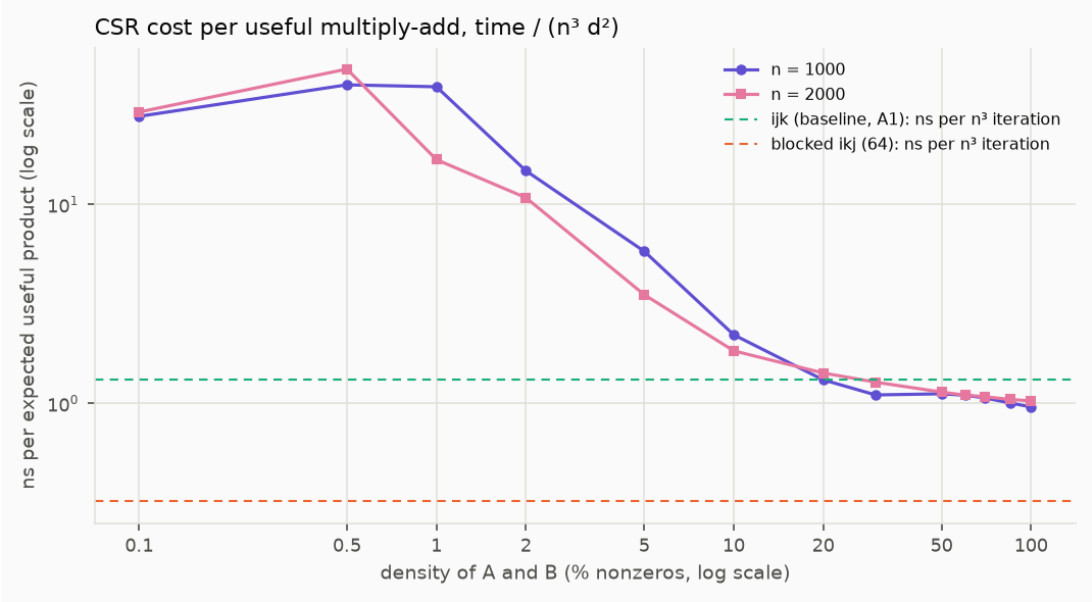


Figure 8. CSR time divided by the expected number of useful products n^3d^2 (d = real density of the inputs), compared with the cost per iteration of the dense methods.

The CSR cost per useful product is 28 to 48 ns at densities of 0.1 and 0.5 %, falls quickly between 1 % and 20 %, and stays at **1.0 to 1.3 ns from 30 % to 100 %** (1.0 to 1.14 ns from 50 %), about 3 to 3.5 times the 0.33 ns per iteration of `blocked`.

4.7 Part B: memory

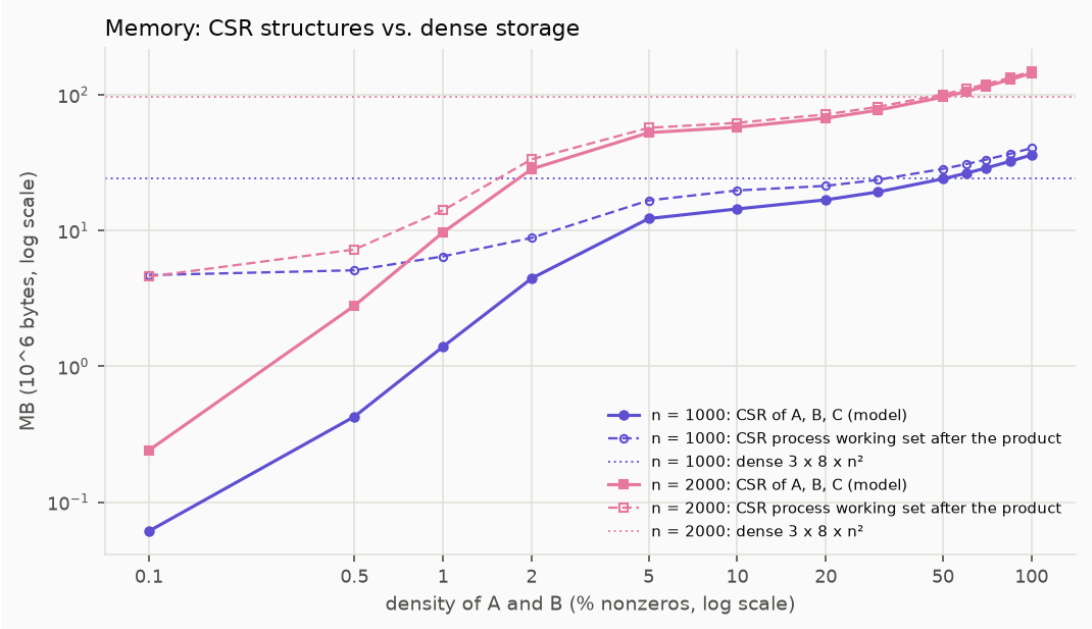


Figure 9. Memory of the CSR structures of A, B and C (model, exact bytes) and working set of the CSR process right after the product, against the dense storage 3 x $8n^2$.

n	density	CSR structures (A + B + C)	CSR working set after product	CSR peak	dense structures	dense peak
1000	0.1 %	0.06 MB	4.7 MB	20.4 MB	24 MB	28.5 MB
1000	5 %	12.3 MB	16.7 MB	34.4 MB	24 MB	28.5 MB
1000	50 %	24.0 MB	28.4 MB	36.5 MB	24 MB	28.5 MB
2000	0.1 %	0.24 MB	4.6 MB	68.5 MB	96 MB	100.5 MB

n	density	CSR structures (A + B + C)	CSR working set after product	CSR peak	dense structures	dense peak
2000	5 %	52.6 MB	57.0 MB	121.3 MB	96 MB	100.5 MB
2000	50 %	96.0 MB	100.4 MB	132.5 MB	96 MB	100.5 MB
2000	100 %	144 MB	148 MB	180.5 MB	96 MB	100.5 MB
4000	0.1 %	1.2 MB	6.1 MB	260.8 MB	384 MB	388.5 MB

The CSR structures are hundreds of times smaller than the dense ones at 0.1 %, but already half the dense size at 5 % (because `C` is almost full), equal at 50 % and 1.5 times larger at 100 %. The working set after the product follows the structures plus about 4.4 MB of the runtime. From 5 % on, the peak of the CSR process is similar to or higher than the dense one (for example 121 MB against 100.5 MB at n = 2000, 5 %), because it includes reading the two dense input files and building the CSR copies.

4.8 Limits

Method	Largest size within the limits	Largest size attempted	Failures
`ijk`	3000 (148 s)	4000	timeout (> 300 s) at n = 4000
`ikj`	4000 (40.0 s)	4000	none
`blocked`	4000 (22.5 s)	4000	none
`csr` (Part B)	4000 at 0.1 % and 1 %	4000	none

No memory failure happened (largest peak 388.5 MB against 4096 MB). `ikj`, `blocked` and `csr` could run larger sizes, but 4000 was the largest declared size; at n = 4000 each dense input file already takes 128 MB.

5. Discussion

5.1 Loop order: the column walk of B was the cause of the Assignment 1 anomalies (H1 confirmed)

Changing only the order of the loops, with the same operations and the same result bit for bit, removes the two memory effects of Assignment 1. In `ijk` the inner loop reads `B` by columns, 8n bytes apart. When `B` ($8n^2$ bytes) no longer fits in the 4 MiB L3 (n of about 720), each read goes to a new cache line from main memory, which explains the jump of `ijk` from 0.74 ns (n = 600) to 1.0 ns (n = 800) and its later growth to 5.5 ns at n = 3000. At n = 512, 1024, 1536 and 2048 the stride is a multiple of 4096 bytes and all the elements of a column fall in the same cache sets, so they evict each other (Section 2): the cost per iteration rises to 5.5 to 16 ns. In `ikj` the inner loop reads rows, so each 64-byte line gives 8 useful values and the stride never matters: the cost stays at 0.53 to 0.73 ns at every size, including the powers of two, and the slope falls from 3.74 to 3.05. This supports, with a controlled experiment, the explanation that Assignment 1 could only propose: the super-cubic growth was a memory-access effect, not a property of the algorithm. The gain grows with n, from 1.0 to 1.6 times up to n = 1200 (except 512 and 1024) to 7.5 times at n = 3000 and 10 to 26 times at the powers of two.

For n <= 50 `ikj` is slower than `ijk`. A plausible reason is that each `(i, k)` pair starts a new inner loop of only n iterations, and for very short loops the loop setup and the final branch weigh more than in `ijk`, whose inner loop accumulates into a register; this was not measured directly.

5.2 Cache blocking: the fastest dense method (H2 partly confirmed)

`blocked` keeps a constant cost of about 0.32 to 0.37 ns per iteration from n = 25 to 4000 (slope 3.04), so it does what H2 predicted for large n. The comparison with `ikj` must be read with care, though: `blocked` is also 1.7 to 1.9 times faster than `ikj` for small n (for example at n = 50, where everything fits in L1 and a single tile

is used). There tiling cannot be the reason, so the speedup of `blocked` over `ikj` in Section 4.2 is not only cache reuse. I did not investigate the cause of this gap at small n (a possible reason is a difference in the machine code that the compiler generated for the two loops), so the cache explanation below is only claimed for large n.

For n <= 500, `B` (at most 2 MB) still fits in the 4 MiB L3, so `ikj` already reuses it from cache; from n = 1000 (`B` = 8 MB), `ikj` must read `B` again from main memory for every row of `C`, while `blocked` reuses each 32 KiB tile of `B` for 64 rows. This fits the gap between the two methods, which is 1.8 to 1.9 times at n >= 1000. The block size results agree with Lam et al. [3]: tiles of 32 x 32 or 64 x 64 (8 or 32 KiB, three of them fit in L2) are best, 128 and 256 (128 KiB and 512 KiB per tile) are 24 to 41 % slower, and 16 x 16 tiles lose some time in the overhead of very short loops. The small rise of `blocked` at n = 2048 (0.43 ns) is probably again a conflict effect: the rows of a tile of `B` are 16 KiB apart, a multiple of the 4 KiB that separates addresses mapped to the same L1 set, so the tile does not fit in L1 as well as at other sizes [3]. The conclusions about the memory effects (Section 5.1) do not depend on the gap at small n, because `ikj` is slower than `blocked` and still removes all the anomalies.

5.3 The baseline at large n is less stable than the other methods

At n = 2000 `ijk` took 15.1 s here and 10.8 s in Assignment 1, and on the sparse inputs of Part B (the same kernel, which does not depend on the values) it varied between 13.5 and 16.1 s depending on the input files. Assignment 1 already reported that the cost of `ijk` above n = 1250 varied without a clear pattern. All the runs here were correct and the IQR at n = 2000 was 5.5 %, so this is not a measurement error but a real sensitivity of the column walk to things that change between processes, probably where the operating system places the physical pages of `B`. `ikj` and `blocked` do not show this behaviour (IQR at n = 2000: 1.3 % and 0.4 %). It also means that the speedups at n >= 1500 have an uncertainty of about +-20 % coming from the baseline.

5.4 Memory of the dense methods (H3 confirmed)

The three dense methods have the same peak memory, equal to 3 x 8 n^2 plus the C runtime, because they use the same arrays and only change the order of the loops. Optimising the access order did not cost memory.

5.5 Sparse multiplication: where CSR helps and where it does not (H4 partly rejected)

H4 predicted that CSR would win at low densities and lose somewhere between 10 % and 50 %. The first part was confirmed, but the crossover is much higher than expected: **about 53 % nonzeros for both n = 1000 and n = 2000**. The cost model of Section 2 explains why. CSR does about n^3d^2 useful products and the dense method does n^3 iterations, so CSR is faster when

```
c_csr x n³ d² < c_dense x n³,   that is   d < sqrt(c_dense / c_csr).
```

With the measured costs at high density, c_csr = 1.0 to 1.14 ns (Figure 8) and c_dense = 0.33 ns (`blocked`), the predicted crossover is sqrt(0.33 / 1.07) = 0.56, very close to the measured 53 %. Both terms grow as n^3, so the crossover does not depend on n, as observed. My hypothesis underestimated how fast the work falls: at 10 % density CSR does only 1 % of the dense products, so even a cost per product three times higher is still a large gain. Each CSR product is about 3 times more expensive than a dense one because of the indirect accesses (`B->col[q]`, then `acc[j]`), the check of `mark[j]` and the scattered writes to the accumulator, while the dense loop reads and writes consecutive positions.

Where the advantage is reduced:

- **Fixed costs at very low density.** At 0.1 % to 0.5 % the CSR cost per useful product is 28 to 48 ns: the work is so small (about 1000 products at n = 1000, d = 0.1 %) that the costs that do not depend on the

products dominate (initialising `mark` and allocating the arrays of size n, the loop over the n rows, sorting or scanning each output row). The time is still tiny in absolute terms (26 µs).

- **Conversion.** Converting the dense input to CSR has to read all n^2 entries, so at low densities it costs far more than the product (2.6 ms against 26 µs at n = 1000, 0.1 %) and reduces the advantage about 100 times. In a real application the data would usually already be stored in a sparse format; with dense input files the conversion should be counted. Its time grows with the density up to 50 to 60 % and then decreases to 100 %; this symmetric shape suggests that the main cost is the branch `if (value != 0)` being mispredicted, since with random positions it is least predictable at 50 %. This is an interpretation; branch mispredictions were not measured.
- **Fill-in of the output.** The density of `C` follows 1 - (1 - d^2)^n: with n = 2000, inputs with 99 % zeros already give an 18 % dense product, and from 5 % the product is practically full. Then the output part of the algorithm (building the CSR of `C` with n^2 entries, sorting the columns) and its memory are as large as for a dense matrix or larger (H5 confirmed).

Memory. The CSR structures take 12 bytes per nonzero instead of 8 bytes per entry. With `C` practically full, the total is about 12 (2d n^2) + 12 n^2, which equals the dense 24 n^2 exactly at d = 50 % (24.0 MB at n = 1000), as measured. So in these experiments CSR saves memory below 50 % and uses more above it, and the saving is limited by the output: at 5 % the structures are already half the dense size although the inputs are 95 % zeros.

The percentage of zeros alone does not determine the result. The same density gives very different advantages for different n (at 1 %: 81 times at n = 1000 and 193 times at n = 2000, product only), because the dense work grows as n^3 and the CSR fixed costs as n or n^2. The output density depends on n d^2, not only on d. Whether the conversion is counted changes the advantage by up to 100 times. And all these results are for uniformly random positions: matrices with structure (banded, block or power-law rows) produce different fill-in and different memory access patterns, which were not tested.

5.6 Limitations

- **One laptop.** The CPU frequency could not be fixed (boost and temperature); new processes, interleaved order, medians and IQR were used, as in Assignment 1. Absolute times and speedups could be different on another machine.
- **No hardware counters.** The cache and branch explanations are based on times and data sizes, not on direct measurements of misses or mispredictions.
- **Sparse inputs.** Only uniformly random patterns were tested, and the inputs are stored in the dense format so that all methods read the same files; this makes the peak memory of the CSR process include the dense inputs.
- **An extra experiment was added after seeing the results** (the high densities, Section 3.9). It is kept separate and labelled as such.

6. Conclusions

These conclusions apply to these implementations (C, `float64`, one thread, gcc 11.2 `-O2 -ffp-contract=off`), these inputs (the Assignment 1 dense files and uniformly random sparse matrices) and this laptop.

1. **Correctness.** The three optimised methods compute exactly the same bits as the Assignment 1 baseline in every case, so the time differences only come from how the same arithmetic is executed.
2. **Loop reordering** (`i-k-j`) removes the memory effects found in Assignment 1: no slowdown at the powers of two and a cost per iteration that stays constant (slope 3.05 against 3.74). It is 1.0 to 1.6 times faster up to n = 1200, 7.5 times at n = 3000 and 10 to 26 times at n = 512 to 2048; it is slower than the baseline only for n <= 50.

3. **Cache blocking** (64 x 64 tiles) keeps the cost at about 0.33 ns per iteration up to n = 4000 and is the fastest dense method (5.7 times faster than the baseline at n = 2000, 15 times at n = 3000). It is 1.7 to 1.9 times faster than `ikj` at every size from n = 100; for n >= 1000 this fits cache reuse (the matrices no longer fit in the L3 cache), but for small n the cause was not investigated. Blocks of 32 or 64 were best.
4. **Dense memory** is the same for the three methods (3 x 8 n^2 plus about 4.5 MB).
5. **CSR (Gustavson)** was faster than the best dense method for all densities up to 50 %, from about 10 000 times at 0.1 % to 1.13 times at 50 %, and slower above about 53 % (3 times slower at 100 %). This crossover was the same for n = 1000 and 2000 and agrees with a simple model: CSR does n^3d^2 products, each about 3 times more expensive than a dense iteration.
6. **The conversion and the output reduce the advantage**: the conversion costs up to 100 times the CSR product at 0.1 %, and the product becomes almost dense from 5 % density, so CSR saves memory only below 50 %.
7. **Limits.** With 300 s per process the baseline reaches n = 3000 (timeout at 4000); `ikj`, `blocked` and `csr` reach the largest declared size, 4000. There were no memory failures.

What remains unresolved.

- Why `blocked` is faster than `ikj` for n <= 500, where the matrices fit in the cache and tiling should not matter.
- The exact cause of the large variation of the baseline at n >= 2000 between processes and between Assignment 1 and 2.
- Whether the explanations based on caches and branch prediction are correct; hardware performance counters would be needed to confirm them.
- How CSR behaves with structured sparse matrices, and how the results change with vectorisation and several threads (Assignment 3).

7. References

[1] Wikipedia, "Matrix multiplication algorithm." [Online]. Available: https://en.wikipedia.org/wiki/Matrix_multiplication_algorithm (accessed Oct. 27, 2026).

[2] Wikipedia, "CPU cache." [Online]. Available: https://en.wikipedia.org/wiki/CPU_cache (accessed Oct. 27, 2026).

[3] M. S. Lam, E. E. Rothberg, and M. E. Wolf, "The cache performance and optimizations of blocked algorithms," in *Proc. 4th Int. Conf. Architectural Support for Programming Languages and Operating Systems (ASPLOS-IV)*, Santa Clara, CA, USA, Apr. 1991, pp. 63-74.

[4] Wikipedia, "Sparse matrix." [Online]. Available: https://en.wikipedia.org/wiki/Sparse_matrix (accessed Oct. 27, 2026).

[5] F. G. Gustavson, "Two fast algorithms for sparse matrices: Multiplication and permuted transposition," *ACM Trans. Math. Softw.*, vol. 4, no. 3, pp. 250-269, Sep. 1978, doi: 10.1145/355791.355796.

[6] 7-cpu.com, "AMD Zen2." [Online]. Available: https://7-cpu.com/cpu/Zen2.html (accessed Oct. 27, 2026).

[7] Free Software Foundation, "Options That Control Optimization," *Using the GNU Compiler Collection (GCC)*, GCC 11.2.0 manual. [Online]. Available: https://gcc.gnu.org/onlinedocs/gcc-11.2.0/gcc/Optimize-Options.html (accessed Oct. 27, 2026).

8. Materials and reproducibility

Everything can be regenerated from the repository https://github.com/shisher06/bigdata-matmul-assignment2. `README.md` has the exact commands to build, run each campaign and make the tables and plots, the versions used and the list of what is reused from Assignment 1 (`bigdata-matmul-assignment1`, commit `e947d18`). For example, Part A is run with:

```
py -3.11 -m pip install -r requirements.txt
py -3.11 scripts/run_benchmarks.py --configs ijk ikj blocked --out results/01_A_densa/raw.csv
```

Item	Location
Source code (four kernels and worker)	`src/c/matmul.c`
Build and flags	`scripts/build.py`, `build/build_info.json`
Input data and configurations	`src/python/datagen.py`, `data/manifest.json` (Assignment 1 files), `data_sparse/manifest.json` (seed, density, nnz and SHA-256 of every sparse matrix)
Tests (24)	`tests/test_workers.py`, `tests/test_data.py`
Benchmark script	`scripts/run_benchmarks.py`
Raw measurements (one row per run)	`results/*/raw.csv`
Environment of each campaign	`results/*/environment.md`
Tables and plots	`scripts/analyze.py`: `results/*/summary.md`, `results/*/plots/`

Each row of the CSV files records the method, block, size, density, numeric type, threads, repetition, batch, warm-up, nonzero counts, conversion time, memory metrics, compiler version and flags. No measurement was removed. Code comments and `results/LEEME.md` are in Spanish.

Use of generative AI. As required by the course policy, the author discloses that a generative AI assistant (Claude Code, Anthropic) was used as a support tool: to help with the code (the new kernels, the CSR code, the changes to the measurement and analysis scripts and the tests), to help interpret results and to help with writing the report. The author is responsible for the content of this report and for defending it.